

MEHREEN HIMANI

Head of AI | AI Governance, Risk & Delivery | Regulated Financial Services

Stuttgart, Germany · +49 152 1516 2753 · himanimehreen@gmail.com · LinkedIn · GitHub

SUMMARY

AI leader with **14 years in regulated financial services** — AML/CFT, credit and market risk, post-trade settlements, cyber security — now building and leading AI capability end to end. I combine the two things most AI leadership roles struggle to hire together: **hands-on fluency with LLM and multi-agent systems** (7 production AI products shipped in the last 12 months, all built by me) and **the regulatory grounding to put them into production inside a bank** (EU AI Act, DORA, GDPR, AML/CFT, Basel III).

I lead AI the way it has to work in a regulated market: use cases triaged by **value x risk**, models selected on **cost per resolved case** rather than benchmark scores, agent autonomy bounded by **guardrails and human-in-the-loop**, and every decision **evidenced in an audit trail a regulator can read**. I have taken AI programmes from a blank page — 12+ use cases prioritised with risk and delivery, LLM enablement rolled out across 4 business units (\$600K+ annual savings identified), a 3x rise in AI pilot proposals within 3 months — and built and mentored distributed teams of 20+ across four time zones to deliver them.

CORE LEADERSHIP CAPABILITIES

AI Strategy & Commercial Ownership — AI operating model design · use-case portfolio and intake funnel · value-vs-risk prioritisation · business case and ROI modelling · build-vs-buy · vendor and model selection · budget, resourcing and delivery economics · proposal and pitch development

AI Governance, Risk & Compliance — EU AI Act (risk classification, Art. 9 risk management, Art. 12 logging, Art. 13 transparency, Art. 14 human oversight, Art. 15 accuracy/robustness) · DORA ICT third-party risk · GDPR (DPIA, Art. 22) · AML/CFT · Basel III · ISO/IEC 42001 and NIST AI RMF alignment · model risk management and AI audit readiness

LLM & Agentic Systems — RAG and hybrid retrieval · multi-agent orchestration and tool design · prompt and context engineering · three-layer guardrails · LLM-as-judge and golden-set evaluation · hallucination and drift monitoring · human-in-the-loop design

AI Cost & FinOps — token economics and cost-per-outcome modelling · model routing and cascades · prompt caching and batch processing · context budgeting · cost guardrails and per-tenant spend caps · unit economics and margin management

Delivery & People Leadership — end-to-end programme delivery in regulated environments · Agile, Scrum, SAgE (PI planning) · onshore/offshore teams of 20+ · hiring, coaching and capability building · C-level and CRO/CFO stakeholder management · multi-region rollout (EMEA, US, APAC)

EXPERIENCE

AI Delivery & Enablement Lead (Consultant) — Capgemini (Sogeti), Stuttgart, Germany

Apr 2023 - Present

AI enablement, quality strategy and programme delivery across banking, capital markets and automotive.

- **Led the LLM-powered AI enablement rollout across 4 business units** end to end — set the adoption roadmap, ran the cross-functional team of consultants and delivery leads, owned Agile cadence, milestones, risks and dependencies. **\$600K+ in annual savings identified.**
- **Built the AI use-case intake and prioritisation model** — ran risk-based triage of 12+ candidate use cases with delivery and risk stakeholders, defining quality gates and delivery milestones per use case. Result: **85% stakeholder alignment and a 3x rise in AI pilot proposals within 3 months** — the practice's demand pipeline went from ad hoc to structured.
- **Designed and stood up the AI evaluation and testing framework** now used to gate LLM features before release — **LLM answer consistency +25%, hallucination rate -40%, model failure rate -30%.**
- **Ship production AI products hands-on** — 7 shipped in 12 months across agentic banking, payments integrity, KYC triage, fraud detection, compliance triage and regulatory RAG (portfolio below). This is what lets me size an engagement, challenge an architecture and earn the trust of engineers.
- Lead the testing and delivery-planning workstream of a **SAP S/4HANA migration** — test strategy, delivery plan, workstream and defect-triage coordination, data migration validation, hands-on system testing.

Project Quality Manager — Standard Chartered Bank, Warsaw, Poland

Mar 2022 - Sept 2022

Credit Risk Analytics and cyber-security programmes; audit risk reduced 40% through integrated governance.

- Owned scope, objectives, deliverables and the full delivery plan for the **Credit Risk Analytics Platform** across **50+ stakeholders** in Risk, Finance and Compliance — validated **Basel III stress-testing and ECB FINREP** workflows, cutting the forecast cycle from **5 days to 8 hours.**
- Delivered the **ETV (Encrypted Traffic Visibility)** cyber-security platform — **AI/ML-driven network traffic analysis** via passive monitoring; owned vendor coordination, execution tracking and risk resolution. Direct experience validating an ML system whose false positives carry operational cost.
- Ran delivery governance across **4 time zones** (Warsaw, London, Mumbai, Singapore) and served as the **escalation point for CRO/CFO** on regulatory deadlines.
- Implemented shift-left testing with BDD and 50+ automated validation checks — **calculation errors -85%, audit findings -60%**; built and mentored a distributed team through hiring, interviewing and coaching.

Project Manager — UBS (via Wipro Technologies), Wrocław, Poland

Dec 2020 – Feb 2022

AML false positives cut 45% through NICE Actimize model testing and tuning.

- Led end-to-end testing and tuning of the **NICE Actimize AML monitoring platform — false positives –45%, 2,000+ analyst hours saved annually, suspicious-activity detection +30%**. The clearest analogue in my background to modern AI risk work: a scored decisioning system, a precision/recall trade-off with real regulatory consequence, and a defensible tuning rationale.
- Owned **estimation, scheduling, resourcing and portfolio invoicing** across onshore/offshore teams; facilitated PI planning for 3 Agile teams — **throughput +25%, cycle time 4 → 2.5 weeks**, 4 major releases on time.
- Ran defect triage with business and compliance stakeholders across US, Swiss and German teams (Corporate Banking, Private Banking, Financial Crime IT).

Project Lead — Credit Suisse (via Wipro Technologies), Wrocław, Poland

Jul 2018 – May 2019

- Led functional and UAT delivery for the **Loan Stress Testing** platform (Private Banking, Zurich); built a quality index that cut **P&L forecasting errors 20%** and drove **30% faster release cycles**.
- Owned delivery and PI planning for 2 Scrum teams — 30+ features over 4 PIs at **90% on-time delivery**.

Test Lead — Credit Suisse (via Wipro Technologies), London, UK / Pune, India

Jan 2013 – Jun 2018

- Led the **multi-region rollout** of the Real-Time Settlements programme across **EMEA, US and APAC** with CI/CD-integrated regression — **regression time –60%, settlement failures –25%**.
- Led a **23-member offshore team** across equity, stock borrow/loan, asset servicing and clearing — **97% SLA adherence, 35% defect-leakage reduction**.
- Delivered market-risk regulatory programmes (**VaR Feed Automation, Volcker, FRTB, TSRD**) with VaR/SVaR validation and daily consolidated reporting.
- **Built “Kandoo,” an internal automation product — \$205K annual savings**, trade testing time –50%, 99.2% coverage with zero critical defects across 8 consecutive releases. My first experience of building an internal product and owning its business case.

Career breaks: Parental leave (Elternzeit) & German language studies — Oct 2022 – Mar 2023; parental leave — Jun 2019 – Nov 2020.

AI PRODUCT PORTFOLIO — BUILT & SHIPPED

Seven production AI products designed, built and shipped in 12 months. Each scoped as a regulated-market product: a named user, a measurable outcome, an evaluation framework and a compliance position.

- **BankFlow AI**: Agentic banking platform — multi-agent orchestration across 8 tools; three-layer guardrails, LLM-as-judge evaluation suite, EU AI Act / AML / DORA compliance embedded in agent reasoning and audit trail.
- **PayClear AI**: Zero-trust payment integrity — pre-execution invoice interception, IBAN verification, visual spoofing detection; composite Trust Index scoring with risk-tiered intervention.
- **OnboardIQ**: Explainable KYC onboarding triage — routes to auto-approve, fast-track or manual review with full decision traceability; a working Art. 14 human-oversight pattern.
- **PayGuard AI**: Real-time payments fraud detection — AI-generated explanations, analyst review queue, EU AI Act compliance score 87/100.
- **ComplianceIQ**: AI compliance alert triage across 6 categories — simulated false positives 45% → 22%, resolution time 8.4h → 4.2h, ~2,100 analyst hours saved annually.
- **PulseAI**: Multi-agent product intelligence — 4-agent pipeline (n8n + Claude API) turning raw feedback into PRDs; ~80% reduction in manual analysis.
- **RegCopilot**: RAG copilot for EU regulatory compliance — 33,536 chunks across EU AI Act, DORA, AMLD6 and GDPR via Supabase pgvector; 82% eval accuracy on a golden set, citation-grounded answers, immutable audit trail satisfying EU AI Act Art. 13.

Flagship reference case study: Project AEGIS — building an enterprise AI function from zero and shipping the first regulated agentic product to production. Covers operating model, EU AI Act conformity mapping, model and vendor strategy, agent architecture, token-level unit economics, risk register and commercial model.

EDUCATION & CERTIFICATIONS

M.Sc. Information Technology — Vellore Institute of Technology, Pune, India · 2013–2017

M.Sc. Biotechnology — Bangalore University, Bangalore, India · 2010–2012

Certifications: AI Agents Certified PM — Product School (Feb 2026) · LLM Foundations Certified PM — Product School (Feb 2026) · Certified Scrum Product Owner® (CSPO®) — Scrum Alliance (Dec 2022) · PMP® Exam Preparation, 35 contact hours

SKILLS

AI & Engineering — GenAI, LLMs, multi-agent systems, RAG, hybrid retrieval, prompt & context engineering, AI evaluation frameworks, LLM-as-judge, guardrail design, human-in-the-loop, AI workflow automation, Claude API, HuggingFace, n8n, Supabase (pgvector), Python, SQL, PostgreSQL, Cursor, Lovable, CI/CD

Governance & Domain — EU AI Act, DORA, GDPR, AML/CFT (NICE Actimize), Basel III, ECB FINREP, credit risk, model risk management, compliance management, enterprise governance, SAP S/4HANA, data migration validation

Leadership & Delivery — AI strategy, operating model design, product roadmapping, OKRs/KPIs, RICE, budget & resource management, vendor management, risk & issue management, Agile/Scrum/SAFe (PI planning), Kanban, Jira, Confluence, ServiceNow, stakeholder management, mentorship, negotiation

Languages — English (C1) · German (B1) · Hindi (native)